

9 Mixed models

9.1 Linear Mixed Models (LMM)

Assumption of the linear model:

- $\epsilon_i \sim N(0, \sigma^2)$
- ϵ_i are **independent** and identically distributed.

Why is independence so important?

Recall the formula for the SE of the mean

$$\frac{\sigma}{\sqrt{n}}$$

Violation of the independence assumption lead to an inflation of Type 1 errors (i.e., rejecting H_0 if it is actually true).

We are often in situations, where we have non-independent data points, for example:

- measurements from the same individual
- measurements from the similar study areas

Possible solutions:

- Include grouping variables in the model as factor.
- Fit individual models to each group
- Use mixed models (fixed and random effects)

9.1.1 Include grouping variables

- If we include a categorical variable, each level (= group) will have its own intercept.
- If we want to have group-specific slopes for an other covariate, we can include interactions.

9.1.2 Fit individual models

- We can fit a model to each group/individual
- Do statistic with coefficients (-> list columns are handy here)

The advantage of this approach is:

- Fewer assumption (no distributional assumption)
- Easier to understand

9.1.3 Use mixed models (fixed and random effects)

We allow individual effects for intercept (and slope) to vary around a common mean.

Recall,

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

9.1.4 Random Intercept Model

$$y_i = \underbrace{(\alpha_j + \beta_0)}_{\text{intercept for group } j} + \beta_1 x_i + \epsilon_i$$

We further assume that

$$\alpha_j \sim N(0, \sigma_\alpha^2)$$

with

- $i = 1, \dots, n$ individuals, and
- $j = 1, \dots, k$ groups.

9.1.5 Random intercept and slope

$$y_i = \underbrace{(\alpha_j + \beta_0)}_{\text{intercept for group } j} + \underbrace{(\gamma_j + \beta_1)x_i}_{\text{slope for group } j} + \epsilon_i$$

We further assume that

$$\begin{bmatrix} \alpha_j \\ \gamma_j \end{bmatrix} \sim N_2(\mathbf{0}, \Sigma)$$

with

- $i = 1, \dots, n$ individuals, and
- $j = 1, \dots, k$ groups, and
- And the variance covariance matrix Σ

$$\Sigma = \begin{bmatrix} \sigma_\alpha^2 & \sigma_{\alpha\gamma} \\ \sigma_{\alpha\gamma} & \sigma_\gamma^2 \end{bmatrix}$$

9.1.6 In R

The package `lme4` allows fitting such models, with the `lmer()` function.

- Random intercept model: $y \sim x + (1 \mid g)$
- Random intercept and slope model: $y \sim x + (x \mid g)$
- Random intercept and slope, and no correlation between slope and intercept (i.e., $\sigma_{\alpha\gamma} = 0$): $y \sim x + (1 \mid g) + (0 + x \mid g)$

9.2 Example

We will use again the data set on mountain pines.

i The `broom` package:

```
library(broom)
```

- The package `broom` provides a set of functions to work with models. The three most used functions are `augment()`, `glance()` and `tidy()`.

```
m1 <- lm(Sepal.Width ~ Sepal.Length, data = iris)
```

The function `augment()` adds several new columns to the data set used to fit the model. These include the fitted (\hat{y}) values of the response (`.fitted`) and the residuals (`.resid`).

```
head(augment(m1))
```

```
# A tibble: 6 x 8
  Sepal.Width Sepal.Length .fitted .resid .hat .sigma .cooksd .std.resid
    <dbl>      <dbl>    <dbl>  <dbl>  <dbl> <dbl>    <dbl>    <dbl>
1         3.5         5.1     3.10  0.397  0.0121  0.435  0.00516     0.919
2          3          4.9     3.12 -0.116  0.0154  0.436  0.000563    -0.269
3         3.2         4.7     3.13  0.0719  0.0195  0.436  0.000277     0.167
4         3.1         4.6     3.13 -0.0343  0.0218  0.436  0.0000709   -0.0798
5         3.6          5     3.11  0.490  0.0136  0.434  0.00893      1.14
6         3.9         5.4     3.08  0.815  0.00859  0.431  0.0154      1.89
```

The function `glance()` provides the main overall information of the model (such as R^2 and adjusted R^2), σ and the values of the F-test as a tibble.

```
glance(m1)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC
    <dbl>      <dbl> <dbl>    <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl>
1  0.0138      0.00716 0.434      2.07  0.152     1 -86.7  179.  188.
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

`tidy()` summarizes a model in terms of the estimated coefficients and returns a *tidy* data.frame.

```
tidy(m1)
```

```
# A tibble: 2 x 5
  term          estimate std.error statistic    p.value
  <chr>        <dbl>    <dbl>    <dbl>    <dbl>
1 (Intercept)    3.42      0.254     13.5 1.55e-27
2 Sepal.Length -0.0619    0.0430     -1.44 1.52e- 1
```

```
dat <- read.table(here::here("data/pines/Data2.txt"), header = TRUE)
```

```
head(dat, 2)
```

	code.tree	site	aspect.tree	NS	EW	z.coord.tree	slope
1	D003	SNP.South.18	214	-0.8290376	-0.5591929	1919.6	34
2	D004	SNP.South.18	221	-0.7547096	-0.6560590	1915.9	30
	tree.height	longest.core.code	longest.core.age	longest.core.dbhbb			
1	10.3	D003.L1	144		14.946		
2	8.3	D004.R1	140		17.258		
	lc.growth.rate50	lc.growth.rate50.cat					
1	0.8400	0.75-1.0 mm					
2	0.9578	0.75-1.0 mm					

```
dat <- data.frame(
  dbh = dat$longest.core.dbhbb,
  gc = dat$lc.growth.rate50.cat,
  age = dat$longest.core.age,
  site = dat$site
)
```

Number of trees per site:

```
table(dat$site)
```


SNP.East.22	SNP.East.24	SNP.East.25	SNP.East.27	SNP.East.28	SNP.North.01
9	9	7	7	8	4
SNP.North.02	SNP.North.03	SNP.North.05	SNP.North.08	SNP.South.11	SNP.South.13
10	9	7	9	6	9
SNP.South.14	SNP.South.18	SNP.South.19	SNP.West.31	SNP.West.32	SNP.West.33
7	7	10	7	8	10
SNP.West.38	SNP.West.39				
9	8				

We start by fitting a simple linear model.


```
m1 <- lm(dbh ~ age, data = dat)
```

- The residuals for each site (the function `add_residuals()` from the `modelr` package is used here to add the residuals to the data.
- If the model is correct, we expect that the site does not have an effect on the residuals.

```
dat %>% add_residuals(m1) %>%
  ggplot(aes(site, resid)) + geom_boxplot() +
  geom_hline(yintercept = 0, col = "red") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```



```
augment(m1) %>% ggplot(aes(age, .fitted)) + geom_line()
```




We could add site as a covariate in the model.

```
m2 <- lm(dbh ~ age + site, data = dat)
```


This allows each site to have its own intercept.

```
augment(m2) %>% ggplot(aes(age, .fitted, col = site)) + geom_line()
```



This improves the residuals

```
dat %>% add_residuals(m2) %>%
  ggplot(aes(site, resid)) + geom_boxplot() +
  geom_hline(yintercept = 0, col = "red") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

But take a look at the number of coefficients that we are fitting now

```
length(coef(m2))
```

```
[1] 21
```

```
nrow(dat)
```

```
[1] 160
```

- Assuming we need 10 - 20 data points for each parameter that we want to estimate, this can be problematic.
- Interpretation becomes tedious at best.

```
coef(m2)
```

(Intercept)	age	siteSNP.East.24	siteSNP.East.25
10.005204952	0.061058802	-4.342966637	1.318461455
siteSNP.East.27	siteSNP.East.28	siteSNP.North.01	siteSNP.North.02
2.295838994	-1.830702836	-0.113083117	-4.431706270
siteSNP.North.03	siteSNP.North.05	siteSNP.North.08	siteSNP.South.11
0.001242395	-4.163689066	-3.833972067	1.059542570

```

siteSNP.South.13 siteSNP.South.14 siteSNP.South.18 siteSNP.South.19
-7.187161295     -5.646999929     -3.128496175     -2.553214615
siteSNP.West.31  siteSNP.West.32  siteSNP.West.33  siteSNP.West.38
-0.431453859    -6.378068692    -6.133800721    -3.467017925
siteSNP.West.39
-6.956053532

```


- The previous model allowed different sites to have different intercepts, but all sites have the same slope.
- We could add an interaction between `age` and `site` to allow for different slopes.

```

m3 <- lm(dbh ~ age * site, data = dat)

augment(m3) %>% ggplot(aes(age, .fitted, col = site)) + geom_line()

```




And again check the residuals, which look even better now.

```

dat %>% add_residuals(m3) %>%
  ggplot(aes(site, resid)) + geom_boxplot() +
  geom_hline(yintercept = 0, col = "red") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



The number of coefficients becomes huge.

```
length(coef(m3))
```

```
[1] 40
```

```
nrow(dat)
```

```
[1] 160
```

An alternative approach would be, to fit for each site the simple model $\text{dbh} \sim \text{age}$.

```
m4 <- dat %>% nest(data = -site)
```

Note, the use of the `nest()` function here. `- nest(data = -site)` groups the data here by site.

```
m4[1:2, ]
```

```
# A tibble: 2 x 2
  site      data
  <chr>    <list>
1 SNP.South.18 <tibble [7 x 3]>
2 SNP.North.03 <tibble [9 x 3]>
```

Finally, a new column `model` (also a list-column) is created with a model for each site.

```
m4 <- m4 %>% mutate(model = map(data, ~ lm(dbh ~ age, data = .)))
```

Each model can now be augmented (entries in the list-column `model`)

```
m4$model[[1]] %>% augment()
```

```
# A tibble: 7 x 8
```

	dbh	age	.fitted	.resid	.hat	.sigma	.cooksd	.std.resid
	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	14.9	144	15.6	-0.647	0.212	2.44	0.0147	-0.331
2	17.3	140	15.4	1.88	0.169	2.24	0.0889	0.936
3	13.0	153	16.1	-3.05	0.387	1.51	0.985	-1.77
4	16.5	141	15.4	1.06	0.178	2.39	0.0305	0.531
5	17.8	134	15.1	2.79	0.143	1.95	0.156	1.37
6	13.1	114	14.0	-0.848	0.398	2.40	0.0815	-0.496
7	12.6	110	13.8	-1.19	0.512	2.31	0.316	-0.776


To this, we have to iterate over `model` list-column and finally `unnest()` the list-column `model`.

```
m4_plot_dat <- m4 %>% mutate(model = map(model, augment)) %>%  
  unnest(cols = model)  
m4_plot_dat[1:2, ]
```

```
# A tibble: 2 x 10
```

	site	data	dbh	age	.fitted	.resid	.hat	.sigma	.cooksd	.std.resid
	<chr>	<list>	<dbl>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	SNP.South~	<tibble>	14.9	144	15.6	-0.647	0.212	2.44	0.0147	-0.331
2	SNP.South~	<tibble>	17.3	140	15.4	1.88	0.169	2.24	0.0889	0.936

```
m4_plot_dat %>% ggplot(aes(age, .fitted, col = site)) + geom_line()
```



- Results look very similar to model m3.
- At some `sites` the relationship between age and DBH is negative.


Now lets use a random intercept model.

```
library(lme4)
library(broom.mixed)
```

- The package `lme4` provides functions to fit mixed effects models.
- The package `broom.mixed` is very similar for the package `broom` (the functions `augment()`, `glance()` and `tidy()`), but for mixed models.

```
m5 <- lmer(dbh ~ age + (1 | site), data = dat)
```

```
m5 %>% augment() %>%
  ggplot(aes(age, .fitted, col = site)) + geom_line()
```




This is very similar to model m2.

```
m2$call
```

```
lm(formula = dbh ~ age + site, data = dat)
```

```
m2 %>% augment() %>%
  ggplot(aes(age, .fitted, col = site)) + geom_line()
```



But compare the number of coefficients between `m2`

```
coef(m2)
```

```
(Intercept)          age  siteSNP.East.24  siteSNP.East.25
10.005204952    0.061058802    -4.342966637    1.318461455
siteSNP.East.27  siteSNP.East.28  siteSNP.North.01  siteSNP.North.02
 2.295838994    -1.830702836    -0.113083117    -4.431706270
siteSNP.North.03  siteSNP.North.05  siteSNP.North.08  siteSNP.South.11
 0.001242395    -4.163689066    -3.833972067    1.059542570
siteSNP.South.13  siteSNP.South.14  siteSNP.South.18  siteSNP.South.19
-7.187161295    -5.646999929    -3.128496175    -2.553214615
siteSNP.West.31   siteSNP.West.32   siteSNP.West.33   siteSNP.West.38
-0.431453859    -6.378068692    -6.133800721    -3.467017925
siteSNP.West.39
-6.956053532
```

```
sigma(m2)
```

```
[1] 3.976816
```

and `m5`

```
fixef(m5)
```

```
(Intercept)          age  
  9.45932393  0.04572627
```

```
VarCorr(m5)
```

```
Groups   Name             Std.Dev.  
site     (Intercept)  2.0533  
Residual                   4.0183
```

In summary

```
# m2  
length(coef(m2)) + 1
```

```
[1] 22
```

```
# m5  
length(fixef(m5)) + 2
```

```
[1] 4
```

Now let us add a random intercept to the existing model.

```
m6a <- lmer(dbh ~ age + (age | site), data = dat)
```

```
Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, :  
Model failed to converge with max|grad| = 1.16634 (tol = 0.002, component 1)
```


```
Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, : Model is near  
- Rescale variables?
```

We get a warning that the optimizer had difficulties to fit the model. The warning message suggest to rescale the covariates (here `age`).


```
m.age <- mean(dat$age)
sd.age <- sd(dat$age)
dat$age1 <- (dat$age - m.age) / sd.age
m6 <- lmer(dbh ~ age1 + (age1 | site), data = dat)
```

After rescaling the covariates, no more warnings are present.

```
m6 %>% augment() %>%
  ggplot(aes(age1, .fitted, col = site)) + geom_line()
```



We observe two things:

- Slopes for different sites differ.
- Slopes are more similar than when fitting individual models. This is called **shrinkage to the mean**.

```
m6
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: dbh ~ age1 + (age1 | site)
Data: dat
REML criterion at convergence: 917.3578
```

Random effects:

Groups	Name	Std.Dev.	Corr
site	(Intercept)	1.934	
	age1	1.059	0.25
Residual		3.954	

Number of obs: 160, groups: site, 20

Fixed Effects:

(Intercept)	age1
16.53	2.45


Correlation between intercept and slope:

```
VarCorr(m6)
```

Groups	Name	Std.Dev.	Corr
site	(Intercept)	1.9338	
	age1	1.0594	0.253
Residual		3.9536	


This means that sites with a higher intercept also grow faster.

```
plot(ranef(m6)$site[, 1], ranef(m6)$site[, 2],  
     xlab = "Intercept", ylab = "age")
```



If no correlation between the slope (`age`) and the intercept is needed, the model would look like this:

```
# Now lets also add a random slope, but no interaction
m7 <- lmer(dbh ~ age1 + (1 | site) + (0 + age1 | site), data = dat)
m7 %>% augment() %>%
  ggplot(aes(age1, .fitted, col = site)) + geom_line()
```



Lets compare the width of confidence intervals:

```
diff(confint(m1)[2, ])
```

```
97.5 %
0.02595968
```

```
diff(confint(m5)[4, ])
```

```
97.5 %
0.03480329
```

```
diff(confint(m6)[6, ])
```

```
97.5 %
2.260624
```

Exercise: Linear Mixed Models

The file `pines/Data1.txt` contains estimates of DBH and age for each tree at different ages. Each tree has an individual id (`core.code`). Perform the following tasks:

1. Plot a growth curve for each tree (in one plot). Where you plot the age on the x-axis and the DBH on the y-axis.
2. Fit the following models:
 - a. A model $DBH = \beta_0 + \beta_1 age$ for each individual tree.
 - b. A global model $DBH = \beta_0 + \beta_1 age$.
 - c. A global model $DBH = \beta_0 + \beta_1 age + \beta_2 id + \beta_3 age \cdot id$.
 - d. The same model as in b), but with a random intercept.
 - e. The same model as in b), but with a random intercept and random slope.
3. Create a plot with the model type (2b, 2d, 2e) on the x-axis and the estimate (with a confidence interval) on the y-axis. Distinguish between the two terms (intercept and slope) by using different panels. Hint, you may find the function(s) `broom::tidy()` with the argument `conf.int = TRUE`, and `bind_rows` useful.

9.3 Some model strategies

9.3.1 How many data levels are needed?

- Literature suggest at least 5 levels are needed.
- Problems can arise if data are spread unequally among levels.
- But see also Oberpriller et al. 2021 [<https://www.biorxiv.org/content/10.1101/2021.05.03.442487v1>].

9.3.2 Choosing random slopes

- Often only random intercepts are used, this can lead to inflated Type 1 error rates (see also Schielzeth and Forstmeier 2009 for details¹).
- It is recommended to fit the most complex model that the data allow.
- See Bates *et al* 2015² for strategies to find the optimal structure of mixed models.

¹<https://academic.oup.com/beheco/article/20/2/416/218997>

²<https://arxiv.org/pdf/1506.04967.pdf>

9.3.3 Model complexity

- Try to avoid fitting the most complex model possible.
- Try to avoid interaction beyond order two.
- You should be able to draw interpretations and visualize your model for different predictors.
- Try to stick to $n/k = 10$, where n is the number of data points and k the number of predictors.

9.3.4 Assessing model fit

- For LMMs residuals analyses as with the linear models are possible.
- Similarly to GLMs we, can use randomized quantile residuals to asses model fit for mixed models with the package **DHARMA**.

9.4 GLMMs

- For generalized linear models with repeated measures and/or hierarchical data, GLMMs can be used. The idea is analogue to LMMs.
- Fitting GLMMs in R can be done with the function `glmer()` from the package **lme4**.
- Randomized quantile residuals can be used to asses model fit.

9.5 Recommended reading

A good place to start is:

- Generalized linear mixed models: a practical guide for ecology and evolution (<https://www.sciencedirect.com/science/article/abs/pii/S0169534709000196>)

More practical guides:

- A brief introduction to mixed effects modelling and multi-model inference in ecology (<https://peerj.com/articles/4794/>)
- Perils and pitfalls of mixed-effects regression models in biology (<https://peerj.com/articles/9522/>)